POST HOCKEY · MODEL CHANGELOG

Model Changelog

Every major model or metric revision, newest first: what changed, why, and what moved.

OCT 29, 2026

All Three Zones: estimates for untracked 2026-27 games are back on

The estimates for games Corey has not tracked yet this season are shown again, from a level correction rebuilt for the refit model. That correction still fails its pre-registered test in one place: for defencemen's carries with a chance, the estimate came in 12.6% under the true count, against a limit of 12.1%. Every other part of the test passed. The limit has not been moved; the estimates are on by an owner's call, and the methods say so.

Why: The miss is small and it is the only one: one stat, for one position, half a percentage point past its line. Without the estimates, the season in progress shows only the games Corey has tracked so far.

What moved: Every skater's 2026-27 All Three Zones number is an estimate for the games played so far again, and the Estimated season view is back for 2026-27. Defencemen's carries with a chance run low, by about an eighth in testing.

Read the writeup

Pro: how each badge's score moves with age

The Pro badge page has an Aging tab: for each skill badge, by forwards and defencemen, how the badge's score changes from one season to the next at each age, corrected for the players who leave the league or lose the role, with a band for what those departed players could change. A curve is drawn only if it beat "no change with age" on a season it was not fitted to and kept its direction when either season was left out; nine of sixteen did, and the rest are listed with the test they failed.

Why: Different badges lose different players, and a curve fitted only to the players who stayed would show the survivors' aging, not the badge's.

What moved: A new tab; no badge, score or holder changes.

Read the writeup

OCT 27, 2026

The defenceman Finisher no longer falls off in October

A defenceman's Finisher badge is his shooting-talent estimate together with the shot value he creates. The talent half was only counted once he had twenty shot attempts in the new season, so on the night it first counted, a week into the season, almost every defenceman fell below the line and seven in ten of last season's holders lost the badge at once. It now needs ten games, like the rest of the badge. On top of that, a defenceman who held the badge at the end of last season keeps it through the early weeks unless he held it narrowly and is in the bottom quarter of the league in shot value after ten games, and a rookie cannot hold it until his score is measured at least as reliably as the least-measured tenth of established defencemen.

Why: The drop was the badge changing its definition overnight, not players starting slowly: in 2024-25 and 2025-26 two thirds of the holders it dropped in October held the badge again by the end of the season.

What moved: From the second week of the season about 34 defencemen hold Finisher instead of 8-12. At season's end it names 3-6 more defencemen a season than before, still about a tenth of them. Forwards are unchanged.

Read the writeup

Playmaker, Transporter, Motor and Speedster ranked against the whole group

These four badges count events: shot and chance assists, carried entries, forecheck recoveries, bursts over 20 mph. A player had to have twenty of the events themselves before he was ranked at all, so only those who already did a lot of the thing were compared, and the top tenth was taken from that slice. Now anyone with ten tracked games is ranked, and the probability carries the thin evidence, as it does for every other badge.

Why: A floor on the count is a floor on the outcome. Since the window weighted the tracking seasons, about five defencemen a season had twenty chance assists, so Playmaker named almost no defencemen, and twice a badge quietly lost one of its two parts (Playmaker's defencemen in 2025-26, Motor in 2025-26).

What moved: Playmaker goes from 0-10 defencemen and 5-6 forwards a season to about 20 and 40; Transporter's defencemen from 2-4 to about 20; Motor from 14-20 forwards in 2023-24 and 2024-25 to about 57, and about 25 of its 2025-26 holders change now that forecheck recoveries count again; Speedster gains about a dozen defencemen and seven forwards a season. Every other badge is unchanged.

Read the writeup

GS+ split into offence and defence

Every GS+ game score is now stored in two halves, offence and defence, that add up to the score exactly. Goals, assists and shots are offence; blocks and shutdown credit are defence; on-ice shot attempts and goals, expected

penalties and faceoffs are each divided between the two, and a goalie's whole score is defence, so his offence half is 0. On the game card the interval view stays the default, and a switch above the card on the share page turns each row into a stacked offence and defence bar. The downloaded image shows whichever view is on screen.

Why: One number for a game says how much, not what kind: a defenceman who earned his score blocking shots and one who earned it shooting read the same until the score is taken apart.

What moved: No GS+ number changed: the formula is the same version (gsplus-3) and the totals match row for row. The halves are new; the card's choice of view is remembered in the reader's browser.

Read the writeup

Lines tab fills in as the game is played

A game's Lines tab no longer waits for the final horn. Each time the poller rescores a game in play, about every minute and a half, the lines are rebuilt from the shift chart as it stands, and the tab says so: "Live: the game so far", with when it was last rebuilt. The league's chart runs a few minutes behind the clock, so the newest minutes arrive late. After the horn they are rebuilt until the shift chart agrees with the boxscore, and the tab reads "the shift chart is still settling" meanwhile; they stay marked provisional until the overnight run writes the permanent lines and replaces them.

Why: Lines are most read while the game is on; waiting for the final horn threw that away.

What moved: A running game shows its lines, built by the same function as the overnight run, which still has the last word. Nothing already shown changes.

Read the writeup

OCT 26, 2026

Lines tab: score & venue adjusted, and filled at the final horn

The Lines tab can show expected goals (sh), shot attempts and goals adjusted for score and venue, the same switch as Linemates; xG impact stays as it was, because its comparison already accounts for score. And a game's lines now fill in about a minute after the final horn, from the poller's own data, marked provisional until the overnight run replaces them.

Why: Lines are most read the night of the game; making readers wait for the overnight run threw away the moment they matter.

What moved: Nothing already shown changes. The overnight run's numbers still replace the provisional ones each morning.

Read the writeup

Skater speed read from every burst, not one top speed

Player pages and the Skaters EDGE tab show hourly burst: the pace a skater reaches about once per 60 minutes of ice time, read from a reversed Gumbel distribution fitted to his season's EDGE burst counts and his ten fastest speeds, not from a single fastest reading. The EDGE tab adds Typical burst, the fit's location, with an 80% range on hover; the spread is one per position, 2.8 mph for forwards and 3.1 mph for defencemen.

Why: A season's single top speed is one reading and swings year to year; a speed read from the whole season's bursts is far steadier.

What moved: New columns and a new chart; EDGE's own max speed is still shown as it was.

Read the writeup

Power-play roles move to team pages, unit by unit

Team pages now show each team's most-used power-play units: where each of

the five skaters played, as shares of the unit's five-on-four time by slot of the 1-3-1, with the five seated in the unit's most likely arrangement. The per-player power-play panel is gone from player pages.

Why: A single role per player claimed more than the model can back; read as a unit, with the split shown, it says what it actually knows.

What moved: The fit behind them is unchanged; its readings are now shown per unit instead of per player.

Read the writeup

OCT 25, 2026

Game pages get a Lines tab: who shared the ice, as units

Each game page has a Lines tab. A team's five-on-five time is cut into forward lines and defence pairs, and shown as a matrix: our lines against our pairs, or our units against theirs. Each cell shows time on ice and one stat at a time: xG impact, shooter-adjusted expected goals, shot attempts or goals. xG impact is the xG (sh) for above what a league-average five would have produced in the same stints, plus the xG (sh) against below it. A unit is a line or pair with at least 10% of its team's five-on-five time, up to six lines and four pairs; everything else is Other, so every row and column adds up and the corner is the team's whole five-on-five time.

Why: A box score lists who played and a shift chart lists who changed; neither says who played with whom, or what those minutes were worth against what an average five would have done there.

What moved: A new tab on game pages, open to everyone; nothing existing changed. One game is very little ice, so the tab is a record of what happened, not a ranking of units.

Read the writeup

Boomerang grades clean exits per retrieval, not retrievals per 60

Boomerang now grades the share of a defenceman's retrieval attempts, clean or botched, that ended in a clean exit, each season against that season's league rate for defencemen, instead of how many retrievals and retrieval exits he made per 60.

Why: A count per 60 rises with how often a defenceman has to go back for the puck, so a team that rarely has to retrieve was penalising its defencemen: the old score correlated 0.89 with his retrieval attempts per 60, the new one 0.04.

What moved: Earned sets change, and most of the holders are different defencemen. Every badge state is replayed under the new definition back to 2023-24, so the change does not show as gains and losses on the front page.

Read the writeup

Experimental: each team's chance of winning every playoff round and the Stanley Cup

Each of the season's simulated seasons now ends in a playoff bracket in the NHL's format, best of seven with home ice, played from the same simulations as the playoff chance. In the playoff games the regular-season strength gap between two teams is cut to 0.48 of itself (c = 0.479), a shrink fitted on 1,048 playoff games from 2013-14 to 2025-26, without 2019-20. Shown in a new Chances view on the standings, on the team page and in Season Predictions, tagged Experimental.

Why: With the full regular-season gap, the stronger team won past playoff series less often than the model said it would; the shrink did better when each past season was scored on a c fitted without it. It is still one fit on twelve postseasons, so it is published as an experiment: the 2026-27 playoffs are its out-of-sample test, and the numbers come down if April contradicts them.

What moved: New columns only. The playoff chance comes out of the same simulations bit for bit, and no other number moves. The front page's model

version changes, because the season simulation's code changed.

Read the writeup

OCT 23, 2026

Badge history recomputed: past seasons' badges moved

Every badge, in every season badges cover (2023-24 on), was recomputed from scratch under the rebuilt models, so who held what in earlier seasons changed, not only this season's. Two rules changed with it. The defencemen's Finisher is now shrunk like any other composite badge; it used to be cut at the empirical 90th percentile, which let an unshrunk shot rate over a couple of dozen games carry a rookie defenceman to it. And where a shrunk score's error comes mostly from a player's own game-to-game spread and rests on fewer than 20 games, he needs a probability of at least 0.50, not the usual line, to earn the badge and to keep it.

Why: Over a dozen games a hot fortnight can clear the line for a call-up or an October rookie, and few badges earned that way were still held in April. A history left on the old models would have set one set of rules beside another.

What moved: Past seasons' badge holders moved. The front page's badge changes are measured under the new rules.

Read the writeup

Zone time over expected, on the Extras tab

Skater Stats' Extras tab has four new even-strength columns. xOZ% and xDZ% are the shares of a skater's ice time we would expect in the offensive and defensive zones, given how his shifts started, whether he is a forward or a defenceman, and his team that season, measured from his teammates with him left out. OZ% vs exp and DZ% vs exp are the gaps between the zone time the NHL's tracking recorded and that expectation, in points, positive for good on

both. One fit over 2021-22 to 2025-26, weighted by minutes; blank under 200 even-strength minutes in a season.

Why: A skater who starts in the offensive zone spends time there; raw zone share partly reports where he was deployed. The gap is what is left once starts, position and teammates are taken out.

What moved: A new column group; nothing existing changed. Unlike the raw share, the gaps follow a player who changes teams (season-to-season r of 0.25 for offensive-zone time and 0.28 for defensive-zone time among players who moved). His linemates' zone time is largely his own, so the gap understates him.

OCT 22, 2026

Every number rebuilt on one revised model

Expected goals: the extra uncertainty added to every shooter's and goalie's talent estimate at the October 2023 opener is gone, the prior on a shooter's talent is tightened by half (variance 0.04 to 0.02), the aging curves behind the talent walk are refit on the new walk, and the base model is solved to convergence with a fixed penalty instead of stopping at an iteration cap. Shift data: zone starts had been misaligned on 30.8% of shift segments since Oct. 9, and are rebuilt for every season; everything that reads them, from implied purpose and shift WPA to QoC, expected giveaways and Game Score+'s shutdown term (re-derived, 7.1 to 7.3), was refit. Skater and team tables, player pages and team pages now reach back to 2010-11. Every season, and this one back to opening night, was recomputed.

Why: The opener's extra uncertainty was meant to let the new tracking data pull talent into place quickly. Held out, it cost accuracy: in the tracking seasons a shooter's talent estimate delivered 0.60 of what it claimed, a goalie's 0.54. The misaligned zone starts fed every number that reads them, and nothing downstream could see it.

What moved: Held-out shot predictions are better in both eras, and every expected-goals figure moved a little, shooter talents toward the average. Everything built on zone starts moved with the correction. The front page's since comparisons were recomputed from the start of the season under the new models, so none of this shows as a move.

Read the writeup

Net Rating computed per 84 games at the source

The model now computes every Net Rating in goals per 84 games itself, instead of computing per 82 and converting on the way to the page, and the Estimated Net Rating and the entrant prior are refit in the same unit.

Why: One unit end to end: a number that is converted once on the way out can be converted twice, or not at all, by the next thing that reads it.

What moved: Net Rating, Estimated Net Rating and their bands were already shown per 84 and do not move because of this. QoC and QoT, the average Net Rating of a shift's opponents and teammates, were still per 82 and are now per 84, about 2.4% larger; their deciles and terciles, and every ranking, are unchanged.

Read the writeup

Game Score+: a goalie keeps his talent and a share of the night

A goalie's GS+ is now split in two: his talent part, what his record says he would save above expected on tonight's shots, and a share of tonight's surprise, the goals he saved or allowed beyond that. The share grows with the shots he faced: a median 23% of it per game over 2023-24 to 2025-26, from 12% on a quiet night to 32% on a busy one. The rest of the surprise is kept in the record and not drawn on the card. Credited goalies are re-centred each season to add up to the league's goals saved above expected. The shutdown weight, re-derived on the rebuilt data, is 7.3, up from 7.1.

Why: One night of saves or goals is mostly shot noise: a goalie's one-game

surprise has a reliability of 0.002. Dom Luszczyszyn's formula pays every save and charges every goal, which over-reads a single night.

What moved: A goalie's number on the game card leans on his record, not on one night's saves.

Read the writeup

All Three Zones: estimates for untracked 2026-27 games are switched off again

The estimates for games Corey has not tracked, switched on for 2026-27 on Oct. 9, are off again. The estimate model was refit on the rebuilt shift data, and the in-season correction that sits on top of it, rerun on the refit, missed its pre-registered test on one cell, defencemen's carries with a chance: a miss of 12.6% against a limit of 12.1%. The page shows Corey's tracked games only for the season in progress until a correction passes. Past seasons keep their estimates.

Why: The Oct. 9 pass was for the old model. A correction that fails its test on the model now in use is not a correction we can publish.

What moved: Every skater's 2026-27 All Three Zones number is again his tracked games only, and the Estimated season view is unavailable for 2026-27.

Read the writeup

OCT 19, 2026

New: power-play roles on every player page

Each skater with 30 minutes of five-on-four in a season now shows how that time split across the five slots of a 1-3-1 power play: net-front, bumper, point and the two half-walls, read from where the unit's events happen, five players at a time. A player with no slot above half is shown as a mixed role.

Why: Who plays where on the power play is the first thing a coach decides about it, and the box score never records it.

What moved: A new panel on player pages; nothing existing changed. The accuracy against hand-labelled units, and the bar the model missed before we chose to publish it, are in the writeup.

Read the writeup

OCT 18, 2026

New: a Standings page with what each team deserved and six ratings

The Standings page has three views. Standings is the NHL table by division, wild card, conference or league, with the last ten games, the streak and the projected points. Deserved asks what the record would be if goals, and the chances behind them, were all that counted: regulation points, Pythagorean and expected-goals winning percentages, and a schedule-adjusted margin. Ratings is six answers to who is number one: Elo, the Simple Rating System, Colley, Keener, Markov and the offense-defense rating, with the mean of their ranks as Consensus. Colley, Keener, Markov and offense-defense are from Langville and Meyer's Who's #1? (Princeton, 2012). Markov is withheld until the median team has played 20 games.

Why: The record says what happened. Whether a team deserved it, and who is actually best, are different questions, and the ratings disagree on purpose: each one decides differently what a win, a margin or a schedule is worth.

What moved: A new page; nothing existing changed.

Read the writeup

Score & venue adjusted on-ice stats, as a toggle on every stats table

Every on-ice shot, goal and xG column gains a score-and-venue adjusted twin, shown in place when the toggle is on, weighted per attempt by the state it was

taken in, with weights refit each season and weekly in season. The talent-controlled weights depart from the textbook McCurdy method at 5v5; see the glossary. Individual counts, the win-probability family and Net Rating are not adjusted.

Read the writeup

OCT 17, 2026

Game Score+ restated in goals

Every GS+ weight is divided by 2.75, so a five-on-five goal with two assists now hands out exactly one goal in all: 0.27 to the scorer, 0.25 and 0.20 to the assists, 0.05 to each of the five skaters on the ice. A power-play goal, or one with fewer assists, hands out less.

Why: In goals, a game's numbers can be read against the score: over 2023-24 to 2025-26 the league's skaters added up to 1.01 times the league's goals.

What moved: Every GS+ reads 2.75 times smaller than before. The division changes no ranking.

Read the writeup

Skater Stats gets a danger filter

Skater Stats' Standard tab can be cut by shot danger, the filter the team page already had: Low, Mid or High, set by each shot's raw expected goals (High is 0.15 or more, Low is under 0.06). The shot, goal and expected-goals columns, on ice and individual, become the chosen bands' sums; Corsi, assists and points, which are not shots in a band, read blank. Ice time and games are the whole selection, so a rate reads as high-danger shots per 60 of that strength. A High danger report comes with it. The danger bands are raw expected goals, so the score-and-venue toggle is off while one is chosen.

Why: Where a team's shots came from says more than how many there were.

The team page had the cut; the players did not.

What moved: A new dropdown; the unfiltered tables are unchanged.

OCT 16, 2026

New: Game Score+ on the game card

The game card ranks each player by Game Score+. It is Dom Luszczyszyn's 2016 Game Score, goals, assists, shots, blocks, penalties, faceoffs and the on-ice shot and goal differentials, each with a fixed weight, with three changes. The faceoff term becomes what a player's draws are worth (JP Gambatese's method). The penalty term becomes the penalties he would be expected to draw and take. And a shutdown term pays a skater for the opponents and the deployment he was handed. Goalies get a line of their own.

Why: The card needs a number that is ready at the final horn, means something after one night, says something about next month and agrees with the three stars, while giving the shutdown defenceman his due. It starts from Dom's Game Score because that came out on top of everything we built or borrowed in our tests.

What moved: A new number on the game card and its share page, with the writeup behind it.

Read the writeup

OCT 14, 2026

Net Rating is now goals per 84 games

From Oct. 14, 2026, Net Rating is shown in goals per 84 games, the length of the NHL season from 2026-27. That covers the season rating, the Estimated Net Rating and its band, the lineup and team totals, the cards and the front page.

Why: The season is 84 games now, and a rating quoted per 82 was a season nobody plays any more.

What moved: Every Net Rating on the site, earlier seasons included, is the old value times 84/82, about 2.4% larger. Rankings, percentiles and badges are unchanged.

Read the writeup

OCT 13, 2026

New: Linemates, any line, pair or group of skaters and what they did together

A new Pro page counts five-on-five time together by stints, stretches of the clock with nobody changing on either side, for forward lines, defence pairs, or any two to five skaters you choose. Each unit shows its time, its zone starts, its quality of competition and teammates, and its expected goals for and against per 60. Expected xGD/60 is what a league-average unit would have produced with the same partners, against the same opponents, from the same starts and score. xG impact (first called xGAx) is the unit's actual shooter-adjusted differential minus that.

Why: Lines are what coaches deploy and the box score never records them. Whether a unit beat what its partners, opponents and starts predicted is a different question from whether it scored. A unit that beats Expected did so in those minutes; the number does not say it will again.

What moved: A new page; nothing existing changed. Each season's page prints the verdict of its pre-registered level checks, a failure included.

Read the writeup

OCT 11, 2026

Full rebuild: every season on the per-second strength split, plus iWPA v2

The second-by-second strength split below now runs through every season, a shift-logging duplicate is merged at the source instead of after the fact, and iWPA is computed under its v2 scheme (13.42).

OCT 9, 2026

5v5 ice time and on-ice stats now split at every strength change

A shift used to carry one strength label for its whole length, even through a penalty taken or expiring mid-shift. Ice time and on-ice stats are now cut second by second at every strength change, and time with a goalie pulled no longer counts as 5-on-5.

Why: A reader compared a player's 5v5 ice time against Natural Stat Trick and the two didn't agree; the strength label on a shift that changed strength partway through was wrong for the part that didn't match how it started.

What moved: Every 2026-27 5v5 number so far — Matvei Michkov's 5v5 ice time, for one, moves from 13:02 to 10:10. Earlier seasons get the same fix this weekend.

FIX

Duplicate shift-chart rows merged

A handful of live-tracked shifts were logged twice, as overlapping rows on the same game's shift chart. They're now merged into one.

What moved: 2026-27 game-page shift charts.

All Three Zones now estimates untracked 2026-27 games

Corey hand-tracks roughly a third of each season; the rest used to show nothing. For games he hasn't tracked yet this season, the page now shows the same kind

of estimate it has long shown for untracked games in past seasons.

Why: Passed a pre-registered test (narrowly), with a per-stat level correction, before turning it on — the same bar every change to this page has to clear.

What moved: Every skater's 2026-27 All Three Zones number that used to read blank.

Read the writeup

OCT. 8–9, 2026

In-game win probability refit

The win-probability model was refit after a model file it depends on changed underneath it, and passed its usual gate before shipping.

Why: Keep the live win-probability chart reading off the same model as the rest of the site, not a stale copy of it.

What moved: Oct. 8's in-game win-probability charts were rebuilt from the stored nightly prior so they match.

OCT 8, 2026

FIX

Badge dates now use the last game played, not the server clock

A badge's "earned since" date is set from the date of the last game that counted toward it, in Eastern time, instead of whatever date the server happened to be on when the badge was computed. Every 2026-27 date on record was re-derived by replaying the season.

OCT 7, 2026

Early-season Net Rating

Before a player has 10 games this season, his Net Rating now blends in last season's position baseline instead of standing entirely on a handful of early games.

Why: A player's whole season shouldn't read off his first few games, especially in the first couple of weeks when nobody has enough of a sample yet.

What moved: Every current-season Net Rating number in October, until each player clears 10 games.

Read the writeup

iWPA shown as a percentage of a win

iWPA now reads as a signed percentage of a win ("+13.1%") everywhere it appears on the site — the game card, the front page, the win-probability panel, Before-and-after — through one shared formatter.

Why: A percentage of a win reads more plainly than the same number stated as a bare fraction.

What moved: Display only; the number underneath hasn't changed, only how it's written.

Read the writeup

Badges score a 250-game window, not the season alone

Every badge is now scored on a skater's last 250 games (about three seasons), each game weighted by how recent it is — a half-life of 55 games for skill badges, 18.3 for role and outcome badges — and carries across the season boundary instead of resetting to nothing every October.

Why: The season-only rule named nobody in the season's first weeks and threw away everything a player had shown the year before. On a three-season holdout, the windowed version predicts the rest of a season better than the season-only one (Pearson r 0.577 against 0.522).

What moved: Badges now carry into a new season instead of starting empty, and a strong start to a season counts toward one sooner.

Read the writeup

OCT 6, 2026

Front-page "since" numbers get one fixed starting point

Front-page movers (playoff chances, Net Rating, and the rest) now compute each comparison's starting value once, overnight, and hold it fixed — rather than comparing against whatever the metric happened to read the last time someone looked.

Why: A moving target makes a trend meaningless; the starting point has to be pinned down once, the same way, for everyone.

What moved: Every "since" callout on the front page.

Model refits recompute front-page “since” comparisons back to the season start, so a trend never shows an artifact of the refit itself. Questions about any entry go to the email on the Glossary page.

Want to help Post Hockey pay for server costs? Consider supporting here.